

MODELOS DE DATA MINING APLICADOS À CLASSIFICAÇÃO DE PULSARES - CÓDIGO

June 1, 2019

1 Classificação de Pulsares usando modelos de Machine Learning

Este trabalho tem por objetivo a Análise da Base de dados HTRU sobre estrelas e pulsares, bem como a modelagem através de modelos de ML e comparação entre diferentes modelos.

2 Importando Bibliotecas e a Base de Dados

```
In [1]: #Importando bibliotecas principais
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # importando ferramentas de pré-processamento e métricas.
```

```
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
#importando modelos
from sklearn.neural_network import MLPClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
```

```
In [3]: #importando dataset e colocando os nomes da colunas
```

```
columns = ['Mean of the integrated profile', 'Standard deviation of the integrated profil
data = pd.read_csv('HTRU_2.csv', names=columns)
```

3 Pré-processamento

Na fase do pré-processamento é feita a análise exploratória do modelo de forma que sejam apagadas inconsistências e valores ausentes nos dados, e que os melhores atributos sejam selecionados para serem usados para predição.

```
In [4]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17898 entries, 0 to 17897
Data columns (total 9 columns):
Mean of the integrated profile      17898 non-null float64
Standard deviation of the integrated profile  17898 non-null float64
Excess kurtosis of the integrated profile  17898 non-null float64
Skewness of the integrated profile    17898 non-null float64
Mean of the DM-SNR curve            17898 non-null float64
Standard deviation of the DM-SNR curve  17898 non-null float64
Excess kurtosis of the DM-SNR curve  17898 non-null float64
Skewness of the DM-SNR curve        17898 non-null float64
Label                               17898 non-null int64
dtypes: float64(8), int64(1)
memory usage: 1.2 MB

```

```
In [5]: data.head()
```

```

Out[5]:  Mean of the integrated profile  \
0      140.562500
1      102.507812
2      103.015625
3      136.750000
4       88.726562

      Standard deviation of the integrated profile  \
0      55.683782
1      58.882430
2      39.341649
3      57.178449
4      40.672225

      Excess kurtosis of the integrated profile  \
0      -0.234571
1       0.465318
2       0.323328
3      -0.068415
4       0.600866

      Skewness of the integrated profile  Mean of the DM-SNR curve  \
0      -0.699648      3.199833
1      -0.515088      1.677258
2       1.051164      3.121237
3      -0.636238      3.642977
4       1.123492      1.178930

      Standard deviation of the DM-SNR curve  \
0      19.110426

```

1	14.860146
2	21.744669
3	20.959280
4	11.468720

	Excess kurtosis of the DM-SNR curve	Skewness of the DM-SNR curve	Label
0	7.975532	74.242225	0
1	10.576487	127.393580	0
2	7.735822	63.171909	0
3	6.896499	53.593661	0
4	14.269573	252.567306	0

In [6]: data.describe()

Out[6]: Mean of the integrated profile \

count	17898.000000
mean	111.079968
std	25.652935
min	5.812500
25%	100.929688
50%	115.078125
75%	127.085938
max	192.617188

Standard deviation of the integrated profile \

count	17898.000000
mean	46.549532
std	6.843189
min	24.772042
25%	42.376018
50%	46.947479
75%	51.023202
max	98.778911

Excess kurtosis of the integrated profile \

count	17898.000000
mean	0.477857
std	1.064040
min	-1.876011
25%	0.027098
50%	0.223240
75%	0.473325
max	8.069522

Skewness of the integrated profile Mean of the DM-SNR curve \

count	17898.000000	17898.000000
mean	1.770279	12.614400
std	6.167913	29.472897

min	-1.791886	0.213211
25%	-0.188572	1.923077
50%	0.198710	2.801839
75%	0.927783	5.464256
max	68.101622	223.392140

	Standard deviation of the DM-SNR curve \
count	17898.000000
mean	26.326515
std	19.470572
min	7.370432
25%	14.437332
50%	18.461316
75%	28.428104
max	110.642211

	Excess kurtosis of the DM-SNR curve	Skewness of the DM-SNR curve \
count	17898.000000	17898.000000
mean	8.303556	104.857709
std	4.506092	106.514540
min	-3.139270	-1.976976
25%	5.781506	34.960504
50%	8.433515	83.064556
75%	10.702959	139.309331
max	34.539844	1191.000837

	Label
count	17898.000000
mean	0.091574
std	0.288432
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
In [7]: sns.heatmap(data.isnull(),yticklabels=False,cbar=False,cmap='viridis')
plt.title('Verificando se existem valores faltantes')
```


```
Out[7]: Text(0.5,1,'Verificando se existem valores faltantes')
```

Verificando se existem valores faltantes



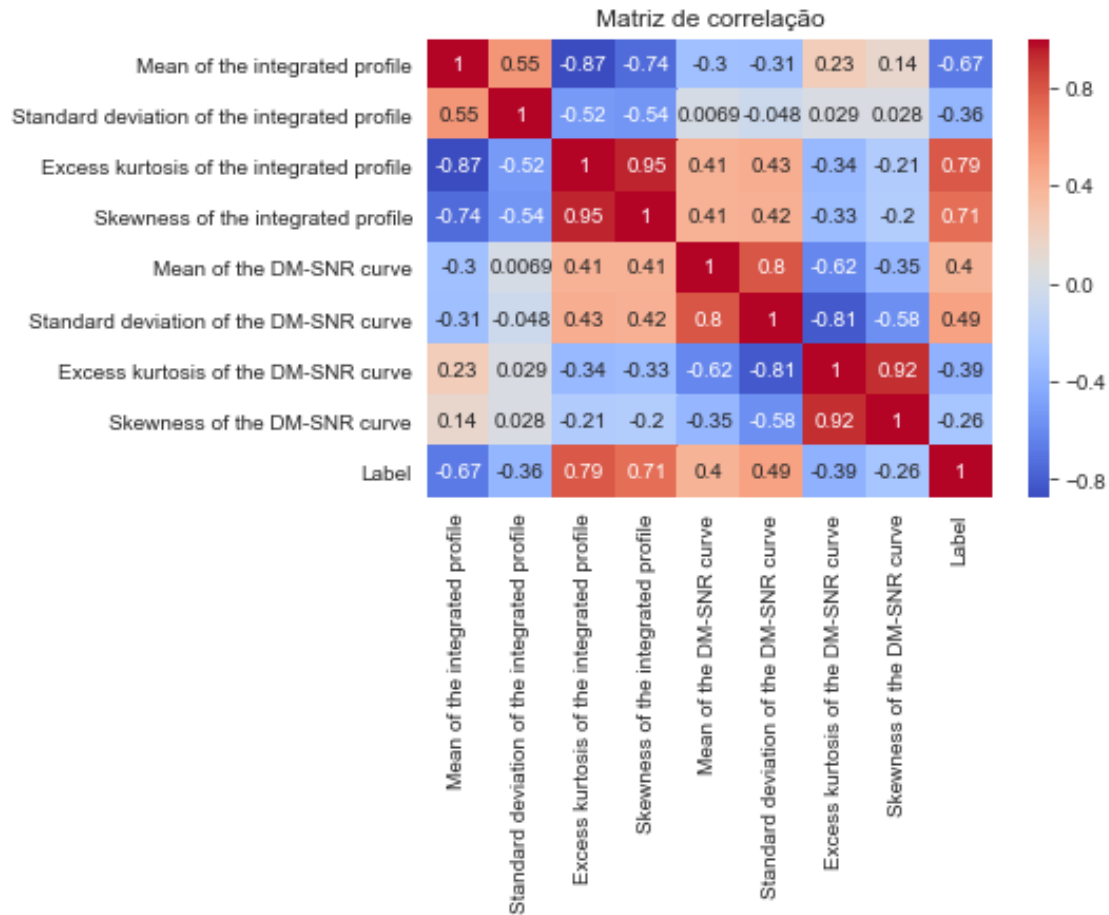
```
In [8]: sns.set_style('whitegrid')
sns.countplot(x='Label',data=data,palette='viridis')
plt.title('Distribuição de Classes')
```

```
Out[8]: Text(0.5,1,'Distribuição de Classes')
```



```
In [9]: sns.heatmap(data.corr(),annot=True, cmap='coolwarm')  
plt.title('Matriz de correlação')  
#Como a base tem 18 mil dados não é necessário excluir nenhum atributo
```

```
Out[9]: Text(0.5,1,'Matriz de correlação')
```



```
In [10]: #instanciando o padronizador dos dados
scaler = StandardScaler()
```

```
In [11]: scaler.fit(data.drop(data.columns[-1],axis=1))
```

```
Out[11]: StandardScaler(copy=True, with_mean=True, with_std=True)
```

```
In [12]: #nova variável de atributos padronizados
scaled_features = scaler.transform(data.drop(data.columns[-1],axis=1))
```

```
In [13]: #Divisão Treino e Teste
X_train, X_test, y_train, y_test = train_test_split(scaled_features,data[data.columns[-1]])
```

4 K Nearest Neighbour (KNN)

```
In [14]: #Escolhendo o melhor K
error_rate = []
```

```

# Vai consumir algum tempo
for i in range(1,42):

    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train,y_train)
    pred_i = knn.predict(X_test)
    error_rate.append(np.mean(pred_i != y_test))

```

```


In [15]: plt.figure(figsize=(10,6))
plt.plot(range(1,42),error_rate,color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')

```

```

Out[15]: Text(0,0.5,'Error Rate')

```



```

In [16]: knn = KNeighborsClassifier(n_neighbors=7)

knn.fit(X_train,y_train)
predictions = knn.predict(X_test)

print('WITH K=7')
print('\n')
print(confusion_matrix(y_test,predictions))

```



```

print('\n')
print(classification_report(y_test,predictions))
print('\n')
knn_score = accuracy_score(y_test,predictions)
print('Acurácia:', knn_score)

```

WITH K=7

```

[[4859  24]
 [ 87 400]]

```


	precision	recall	f1-score	support
0	0.98	1.00	0.99	4883
1	0.94	0.82	0.88	487
avg / total	0.98	0.98	0.98	5370

Acurácia: 0.9793296089385475

```

In [17]: plt.figure(figsize=(10,7))
sns.heatmap(confusion_matrix(y_test, predictions), annot=True, fmt='.1f', cmap=sns.col
plt.title('Matriz de confusão KNN')
plt.show()

```



5 Gaussian Naive Bayes (GNB)

```
In [18]: gnb = GaussianNB()

In [19]: gnb.fit(X_train,y_train)
         predictions = gnb.predict(X_test)

In [20]: print(confusion_matrix(y_test,predictions))
         print('\n')
         print(classification_report(y_test,predictions))
         print('\n')
         gnb_score = accuracy_score(y_test,predictions)
         print('Acurácia:', gnb_score)

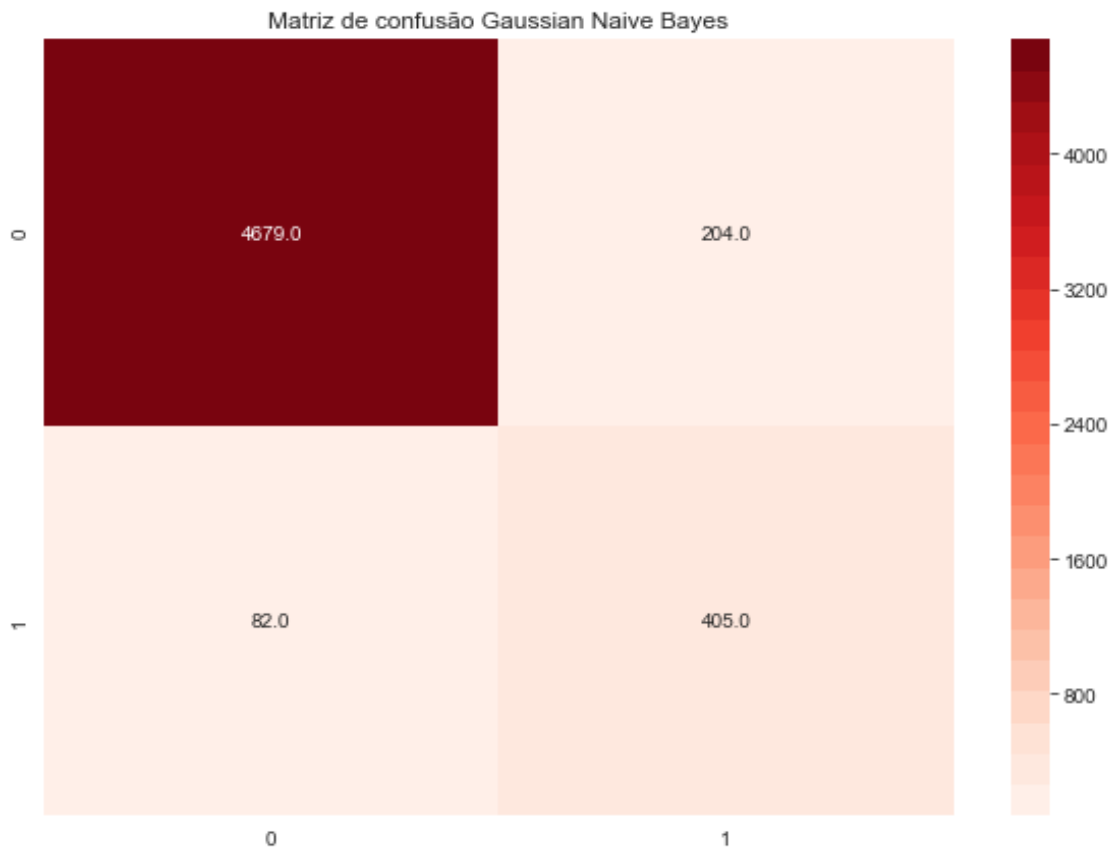
[[4679  204]
 [   82 405]]
```

```
precision    recall  f1-score   support
```

0	0.98	0.96	0.97	4883
1	0.67	0.83	0.74	487
avg / total	0.95	0.95	0.95	5370

Acurácia: 0.9467411545623836

```
In [21]: plt.figure(figsize=(10,7))
sns.heatmap(confusion_matrix(y_test, predictions), annot=True, fmt='.1f', cmap=sns.colorm
plt.title('Matriz de confusão Gaussian Naive Bayes')
plt.show()
```



6 Redes Neurais Artificiais (RNA)

Rede com uma camada oculta

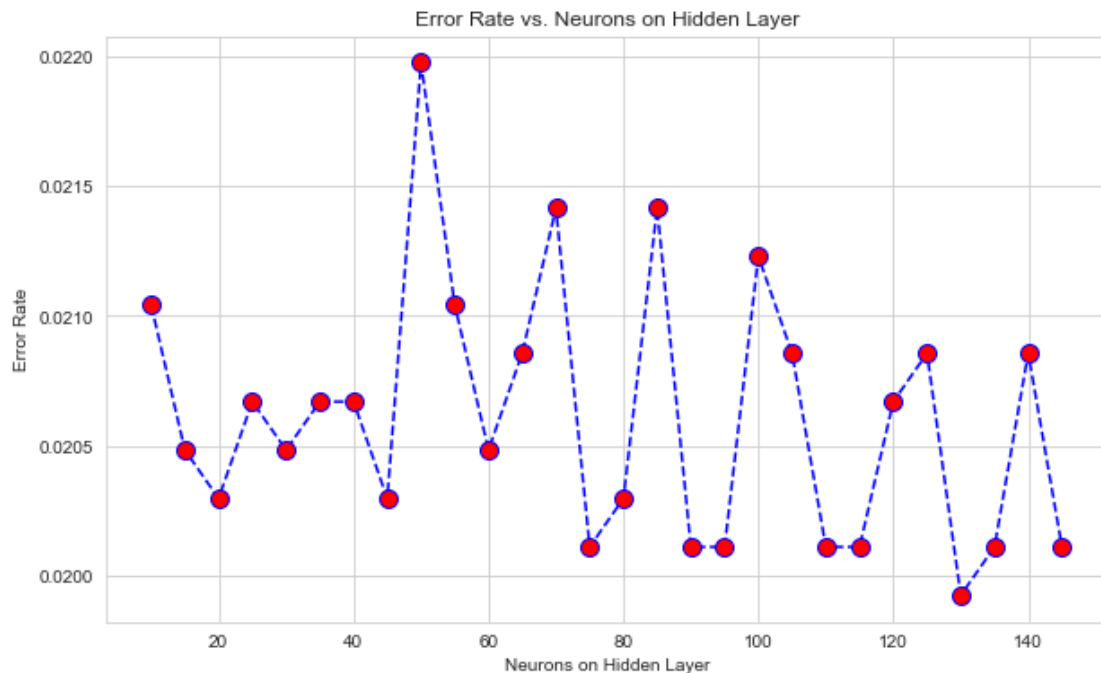
```
In [24]: error_rate = []
```

```
# Vai demorar alguns minutos
for i in range(10,150,5):

    mlp = MLPClassifier(hidden_layer_sizes=(i, ), activation='logistic')
    mlp.fit(X_train,y_train)
    pred_i = mlp.predict(X_test)
    error_rate.append(np.mean(pred_i != y_test))
```

```
In [25]: plt.figure(figsize=(10,6))
plt.plot(range(10,150,5),error_rate,color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)
plt.title('Error Rate vs. Neurons on Hidden Layer')
plt.xlabel('Neurons on Hidden Layer')
plt.ylabel('Error Rate')
```

```
Out[25]: Text(0,0.5,'Error Rate')
```



```
In [26]: mlp = MLPClassifier(hidden_layer_sizes=(75, ), activation='logistic')
mlp.fit(X_train,y_train)
predictions = mlp.predict(X_test)
```

```
In [27]: print('WITH 75 Neurons on Hidden Layer')
print('\n')
print(confusion_matrix(y_test,predictions))
print('\n')
print(classification_report(y_test,predictions))
```

```

print('\n')
mlp_score = accuracy_score(y_test,predictions)
print('Acurácia:', mlp_score)

```

WITH 75 Neurons on Hidden Layer

```

[[4858  25]
 [  83 404]]

```

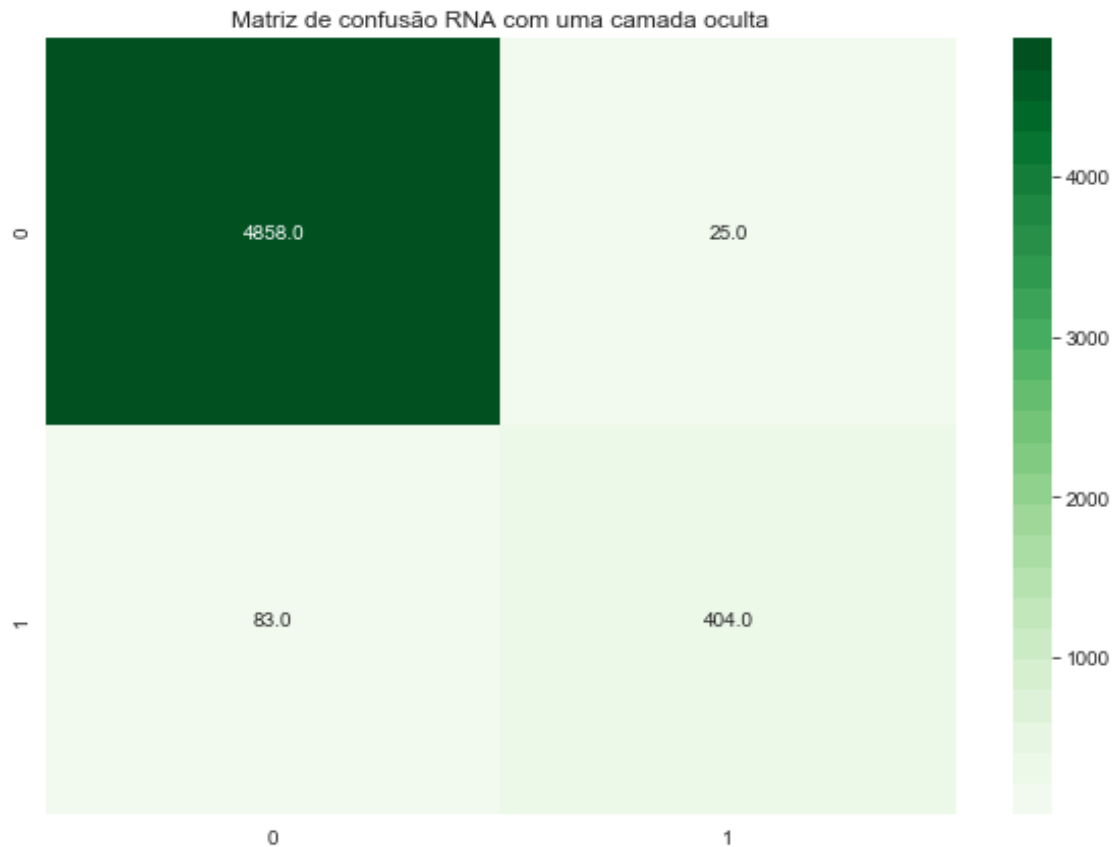
	precision	recall	f1-score	support
0	0.98	0.99	0.99	4883
1	0.94	0.83	0.88	487
avg / total	0.98	0.98	0.98	5370

Acurácia: 0.9798882681564246

```

In [28]: plt.figure(figsize=(10,7))
sns.heatmap(confusion_matrix(y_test, predictions), annot=True, fmt='.1f', cmap=sns.col
plt.title('Matriz de confusão RNA com uma camada oculta')
plt.show()

```



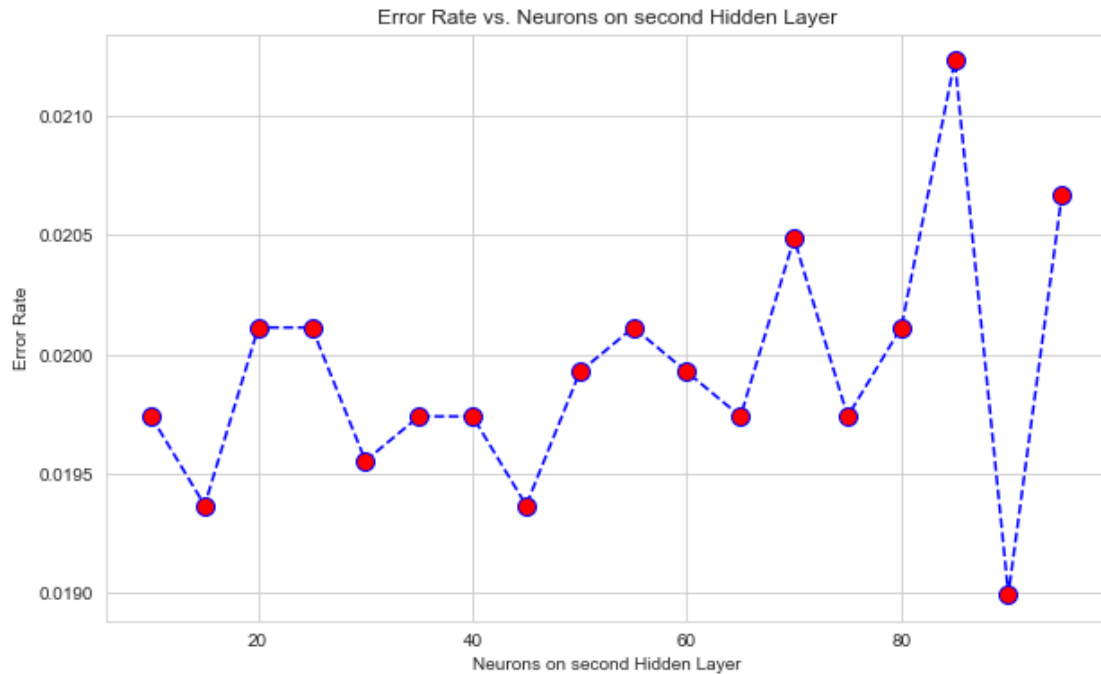
Verificando o ganho de desempenho com a adição de uma nova camada oculta

```
In [31]: # Vai demorar alguns minutos
error_rate = []
for i in range(10,100,5):

    mlp = MLPClassifier(hidden_layer_sizes=(75,i, ), activation='logistic')
    mlp.fit(X_train,y_train)
    pred_i = mlp.predict(X_test)
    error_rate.append(np.mean(pred_i != y_test))

In [32]: plt.figure(figsize=(10,6))
plt.plot(range(10,100,5),error_rate,color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)
plt.title('Error Rate vs. Neurons on second Hidden Layer')
plt.xlabel('Neurons on second Hidden Layer')
plt.ylabel('Error Rate')

Out[32]: Text(0,0.5,'Error Rate')
```



```
In [33]: mlp = MLPClassifier(hidden_layer_sizes=(75, 45, ), activation='logistic')
mlp.fit(X_train,y_train)
predictions = mlp.predict(X_test)
```

```
In [34]: print('WITH 75 Neurons on Second Hidden Layer')
print('\n')
print(confusion_matrix(y_test,predictions))
print('\n')
print(classification_report(y_test,predictions))
print('\n')
second_mlp_score = accuracy_score(y_test,predictions)
print('Acurácia:', second_mlp_score)
```

WITH 75 Neurons on Second Hidden Layer

```
[[4860  23]
 [  84 403]]
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	4883
1	0.95	0.83	0.88	487

avg / total	0.98	0.98	0.98	5370
-------------	------	------	------	------

Acurácia: 0.9800744878957169

```
In [35]: plt.figure(figsize=(10,7))
sns.heatmap(confusion_matrix(y_test, predictions), annot=True, fmt='.1f', cmap=sns.col
plt.title('Matriz de confusão RNA com duas camadas ocultas')
plt.show()
```

